



# Module 4 – Part I

## Machine Learning Fundamentals


---





# Class Modules

- Module 1- Introduction to Deep Learning
- Module 2- Setting up Machine Learning Environment
- Module 3- Linear Regression (Econometrics approach)
- **Module 4- Machine Learning Fundamentals**
- Module 5- Linear Regression (Machine Learning approach)
- Module 6- Penalized Regression (Ridge, LASSO, Elastic Net)
- Module 7- Logistic Regression
- Module 8- K-Nearest Neighbors (KNN)
- Module 9- Classification and Regression Trees (CART)
- Module 10- Bagging and Boosting
- Module 11- Dimensionality Reduction (PCA)
- Module 12- Clustering (KMeans – Hierarchical)







# Topics

---

- Inference vs Prediction
- The Model
- Train, Test, Validation
- Resampling methods
- Evaluation metrics
- Bias-Variance tradeoff, overfitting, learning curve
- How do machines actually learn?
  - Cost Function
  - Solvers/learners (GD, SGD)
- Scaling the features



# The Model



# The Model

---

$$y = f(X, \theta) + \epsilon = f(X_1, X_2, \dots, X_m, \theta_1, \theta_2, \dots, \theta_k) + \epsilon$$

$y$  : response, dependent variables, output, **Target**

$X$ : predictors, independent variables, input, **Features**

$\theta$ : estimates, specifications, **Parameters**

✓ It is all about estimating  $f$  by  $\hat{f}$  for two purposes:

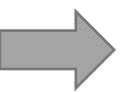
- 1) Inference (interpretable ML)
- 2) Prediction

# → Parameters and Hyperparameters

$$y = f(X, \theta) + \epsilon = f(X_1, X_2, \dots, X_m, \theta_1, \theta_2, \dots, \theta_k) + \epsilon$$

Model **parameters** are estimated from data automatically and model **hyperparameters** are set manually (prior to training the model) and are used in processes to help estimate model parameters.

Example?



# Parametric Vs. Nonparametric models

$$y = f(X, \theta) + \epsilon$$

- The true relationship,  $f(X)$  is unknown and the goal is to see which ML algorithm is better at approximating it. An algorithm learns/estimates  $f(X)$  from training data.

$f(X)$  is **assumed**. Examples:  
Linear regression, GLM,  
logistic regression, simple  
Neural networks, ....



Parametric  
algorithms

Pros



**Simpler**

Easier to understand and to interpret

**Faster**

Very fast to fit your data

**Less data**

Require "few" data to yield good perf.

Cons



**Limited complexity**

Because of the specified form, parametric algorithms are more suited for "simple" problems where you can guess the structure in the data

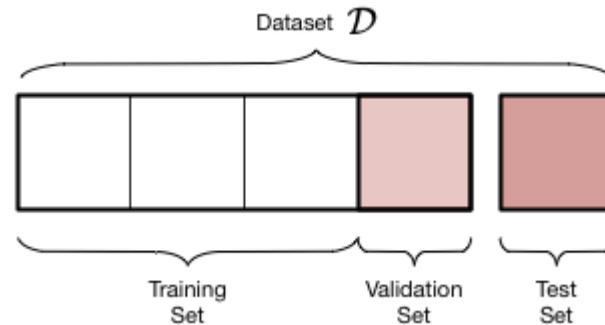
# Train, Validation, Test Resampling Methods



# → Partitioning of the dataset

The data set is typically divided into three non-overlapping samples:

- 1) **Training set:** to train the model
- 2) **Validation set:** to validate and tune the model
- 3) **Test set:** to test the model's ability to predict well on new data (**generalize**)



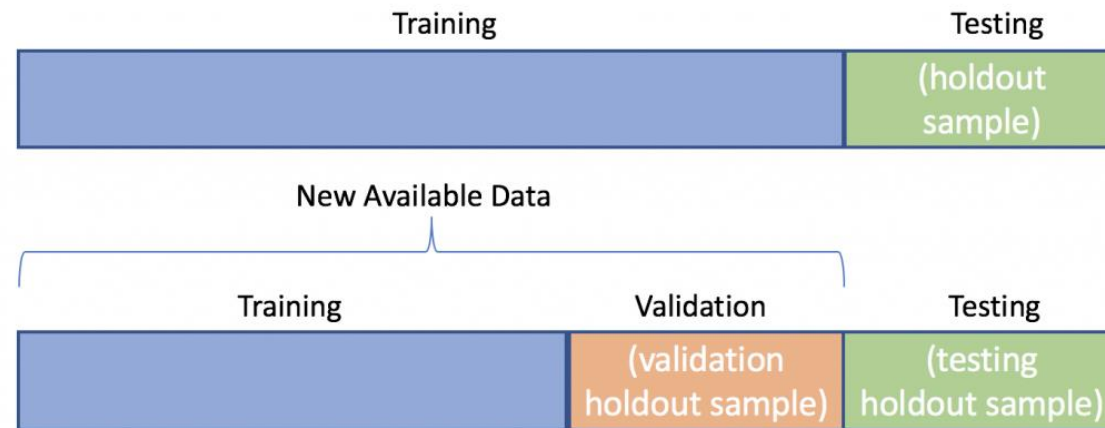
To be valid and useful, any supervised machine learning model **must** generalize well beyond the training data.

**Large dataset is needed! But what if we don't have it?**



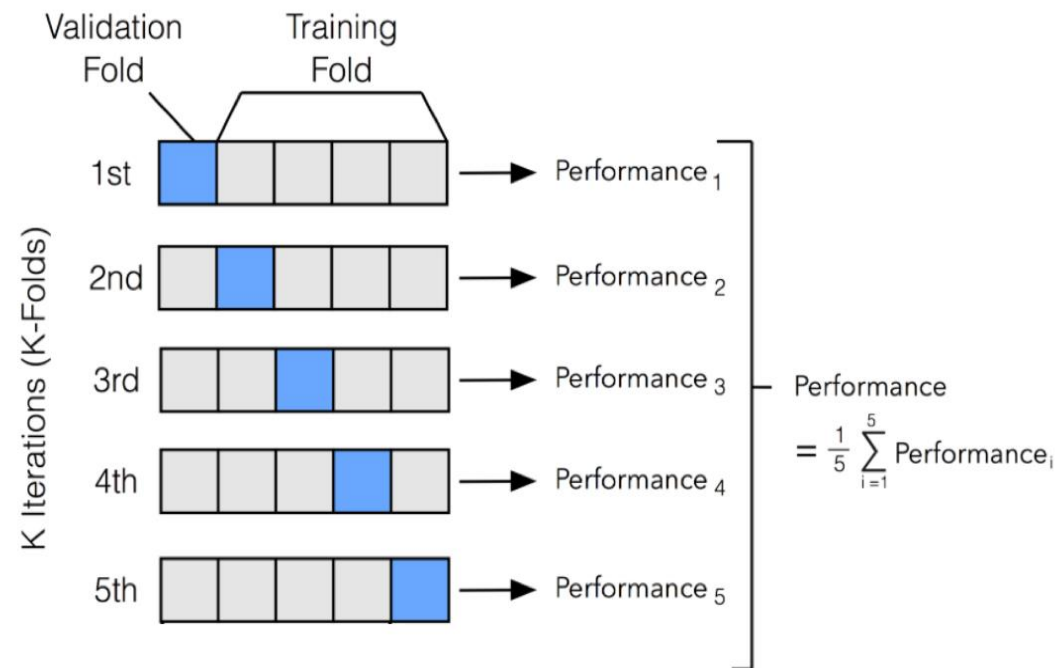
# Resampling methods

- Sometimes we cannot afford to split the data in three because the algorithm may **not learn** anything from a **small training dataset**!
- **Small validation set** is also problematic because we cannot tune the hyperparameters properly! **Unstable** model performance in validation set!
- **Solution**: combining the training and validation sets and use cross validation!



# ➔ K-fold Cross Validation

- 1) Divide the training data into  $K$  roughly equal-sized non-overlapping groups. Leave out  $k^{th}$  fold and fit the model to the other  $k - 1$  folds. Finally, obtain predictions for the left-out  $k^{th}$  fold.
- 2) This is done in turn for each part  $k$  and then the results are combined.



# ➔ Why do we use Cross Validation?

Cross validation is mainly used for two purposes:

1. Model **architecture** selection (optimization vs generalization)
2. **Estimation** of model performance in the test set



- After selecting the **best model architecture**, we estimate the generalization error using the test set.
- **Different model comparison** is based on **test set** performance!

# Evaluation Metrics




# Evaluation metrics

In general, we want to compare how close are the predictions to the actual numbers in the **test set**.


This is typically assessed using

- MSE for **quantitative** response
- Misclassification rate for **qualitative** response

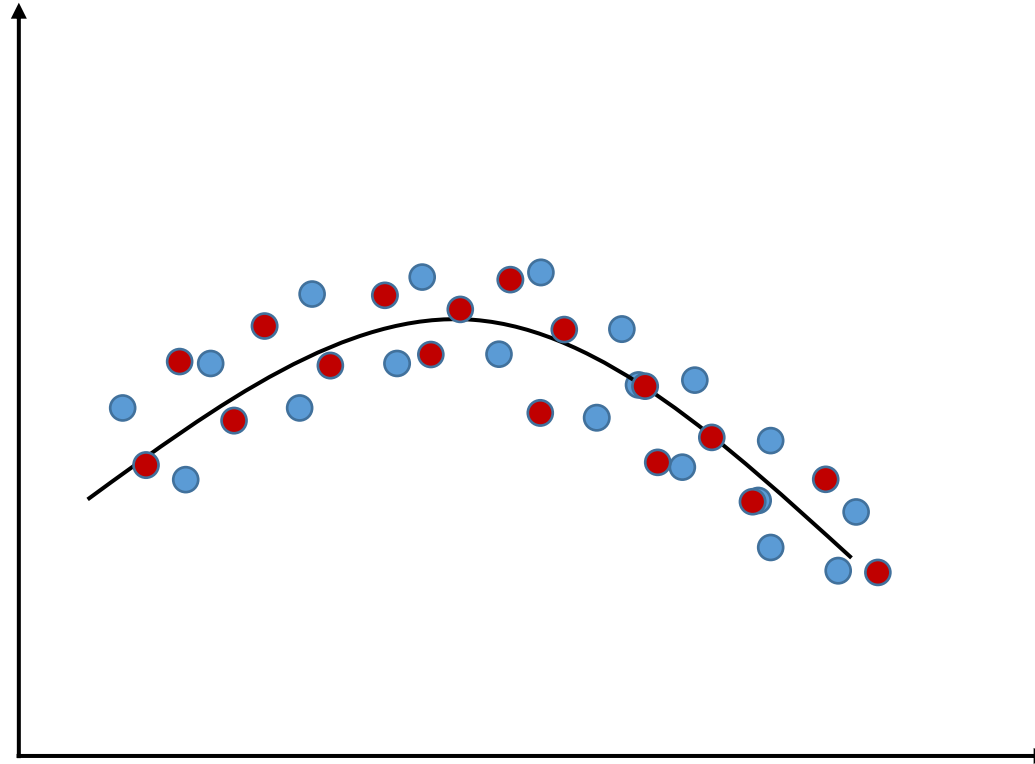


# Bias-Variance Tradeoff

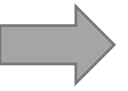
## Optimization vs Generalization



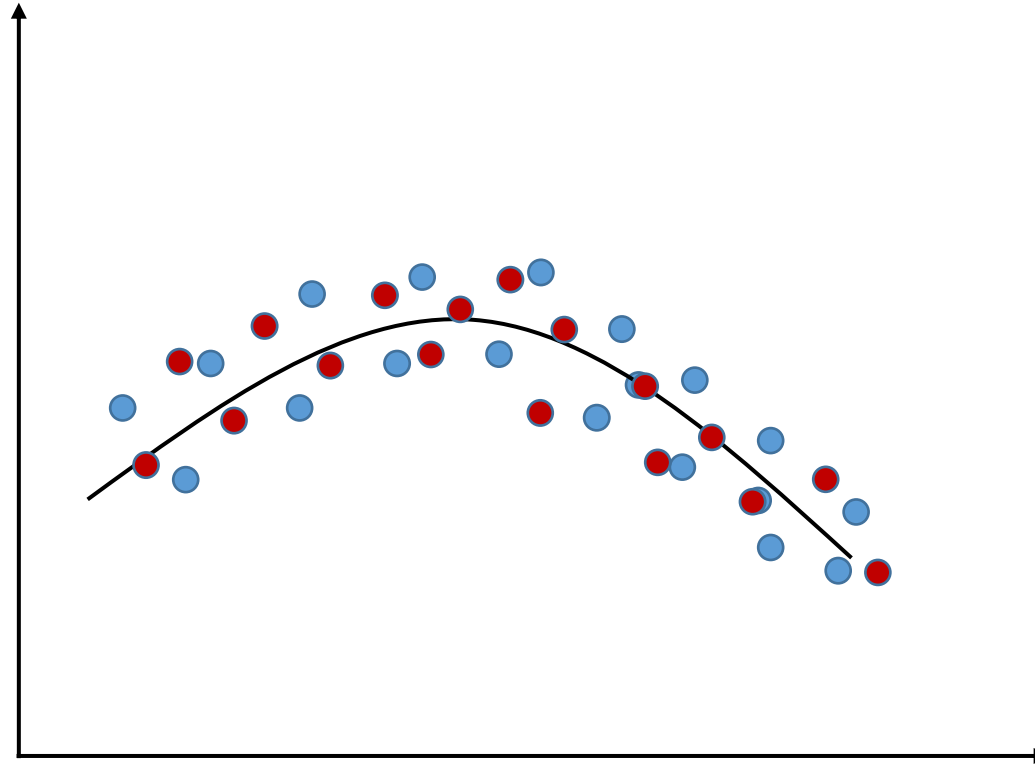
# Model Bias & Model Variance in machine learning







# Model Bias & Model Variance in machine learning



# → MSE decomposition

$$MSE = \text{model variance} + \text{model bias} + \text{irreducible error}$$

- 1) **Model variance** is the variance if we had estimated the model with a different **training set**
  - 2) **Model bias** is the error due to using an approximate model (model is too simple)
  - 3) **Irreducible error** is due to missing variables and limited samples. Can't be fixed with modeling
- The goal is to minimize the sum of **model variance** and **model bias**.
  - This is known as the bias-variance tradeoff because reducing one often leads to increasing the other.
  - Choosing the flexibility (complexity) of  $\hat{f}(X)$ , will amount to bias-variance tradeoff.

# ➔ MSE decomposition

The **bias-variance** tradeoff is one of the core concepts in supervised learning.



Assume that the data is generated by a simple model!


$$y_i = f(\mathbf{x}_i) + \epsilon_i, \quad \mathbb{E}[\epsilon] = 0, \quad \mathbb{V}[\epsilon] = \sigma^2$$

The estimated model yields



$$\hat{y}_i = \hat{f}(X_i)$$

Let us decompose the mean squared error (**MSE**):

$$\begin{aligned} \mathbb{E}[\hat{\epsilon}^2] &= \mathbb{E}[(y - \hat{f}(\mathbf{x}))^2] = \mathbb{E}[(f(\mathbf{x}) + \epsilon - \hat{f}(\mathbf{x}))^2] \quad \dots = \underbrace{\mathbb{V}[\hat{f}(\mathbf{x})]}_{\text{variance of model}} + \underbrace{\mathbb{E}[(f(\mathbf{x}) - \hat{f}(\mathbf{x}))^2]}_{\text{squared bias}} + \sigma^2 \\ &= \underbrace{\mathbb{E}[(f(\mathbf{x}) - \hat{f}(\mathbf{x}))^2]}_{\text{total quadratic error}} + \underbrace{\mathbb{E}[\epsilon^2]}_{\text{irreducible error}} \end{aligned}$$




# Representations of the bias-variance tradeoff

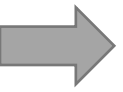


Optimization  
Vs  
Generalization



# Other representations of the bias-variance tradeoff











# Overfitting

Overfitting happens when the fitted algorithm does **not generalize** well to new data:



- The model fits the training data **too** well while not predicts well in the new data
- The model **fits the noise** ( $\epsilon$ ) in training data (finds a pattern that does not exist)
- The algorithm has simply **memorized** the data, rather than **learned** from it!
- The model is too **complex**!



# ➔ Mitigate overfitting

The main techniques used to mitigate overfitting risk in a model construction are:

- 1) Collect **more data** (Can reduce AND variance)
- 2) **Complexity** reduction (regularization, feature selection)
- 3) **Cross validation** (estimate the performance in test set)




With more training example



# The Learning Curve: Do we need to collect more data?

- A learning curve is a plot that shows the relationship between the amount of **training data** and the **performance** of a machine learning model.
- It is used to diagnose whether a model has high **bias**, high **variance**, or is **just right**.


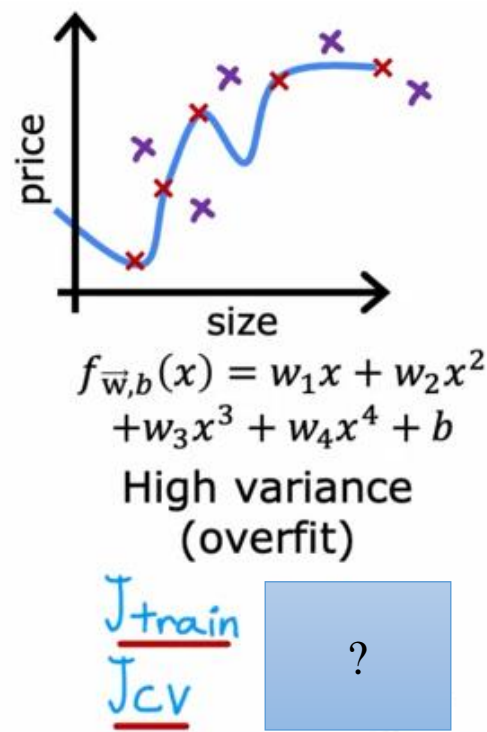

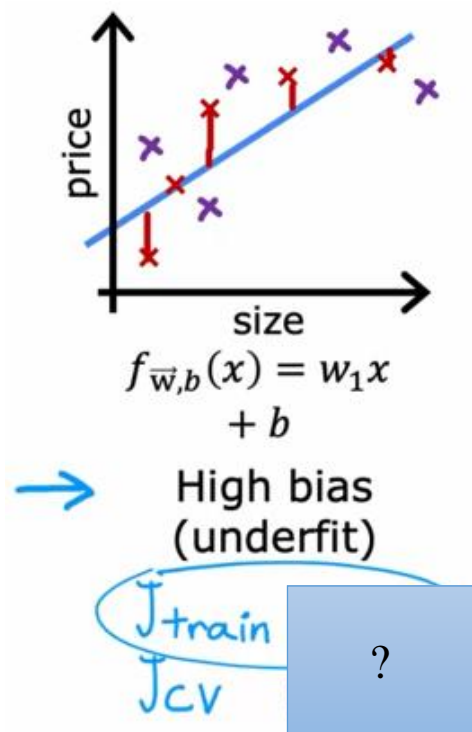
-- Training score  
-- Cross validation score  
-- Benchmark performance  
(common sense performance)







# Class exercise




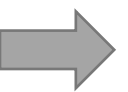
# How do Machines Learn?



# Terminology

- Learning: Finding the model **weights** (parameters' values)
- **Cost Function**: Tells us “**how good**” our model is at making predictions for a given set of parameters.
- The cost function has its own curve and its own gradients. The slope of this curve tells us how to update our parameters to make the model more accurate.
- The two most frequently used optimization algorithms when the cost function is **continuous** and **differentiable** are Gradient Descent (GD) and Stochastic GD.





# Solvers (learners)!


- **Gradient Descent**: is an **iterative** optimization algorithm for finding the minimum of a function.
- We start at some random point and take steps proportional to the **negative** of the gradient of the function at the current point.

$$\theta_j := \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta)$$

- $\theta_j$  is the model's  $j^{th}$  parameter
- $\alpha$  is the learning rate
- $J(\theta)$  is the cost function (which is differentiable)

# Choice of learning rate

- If  $\alpha$  is too small, gradient descent can be slow
- If  $\alpha$  is too large, the gradient descent can even diverge.

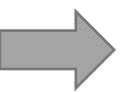


# ➔ Beyond Gradient Descent?


**Disadvantages** of gradient descent:

- Single batch: use the entire training set to update parameters!
- Sensitive to the choice of the learning rate
- Slow for large datasets


**(Minibatch) Stochastic Gradient Descent**: is a version of the algorithm that speeds up the computation by approximating the gradient using **smaller batches** (subsets) of the training data. SGD itself has various “upgrades”.




# SGD vs GD




- Batch gradient descent
- Mini-batch gradient Descent
- Stochastic gradient descent





# Beyond SGD?


- Loss functions can be difficult to optimize!
- *Visualizing the loss landscape of neural nets*, Li et al, 2018



- **Solution:** Designing an adaptive learning rate that can **adapt** to the loss landscape.
- Rather than just looking at **the current gradient**, consider the **previous weight updates**.
- This is called, **momentum**!
- Examples: Adam, Adadelata, Adagrad, RMSProp!



# How do machines learn?




# Feature Scaling


# ➔ Why feature scaling?

- Feature scaling in machine learning is a critical step during the **pre-processing** of data **before creating a machine learning model**.
- Feature scaling is essential for machine learning models that calculate **distances between data**.
- Feature scaling could:
  - Avoid numerical overflow and speed up the algo
  - Reduce dominant effects of specific variables

$$\begin{aligned} 0 \leq x_1 \leq 1 \\ 0 \leq x_2 \leq 1 \end{aligned}$$



Both parameters could be updated in equal proportions




Gradient of larger parameters dominates the updates

# ➔ Scaling the features




Let us use  $x_i$  for raw input and  $\tilde{x}_i$  for the transformed data.  
Common scaling practices include:

- Standardization (Z-score):
$$\tilde{x}_i = \left( \frac{x_i - \mu_x}{\sigma_x} \right)$$
- Normalization:
  - Min-Max scaler over [0,1]:
$$\tilde{x}_i = \left( \frac{x_i - \text{Min}(X)}{\text{Max}(X) - \text{Min}(X)} \right)$$
  - Min-Max scaler over [-1,1]:
$$\tilde{x}_i = 2 * \left( \frac{x_i - \text{Min}(X)}{\text{Max}(X) - \text{Min}(X)} \right) - 1$$
  - Mean normalization:
$$\tilde{x}_i = \left( \frac{x_i - \text{Mean}(X)}{\text{Max}(X) - \text{Min}(X)} \right)$$





# Scaling the features (class exercise)



aim for about  $-1 \leq x_j \leq 1$  for each feature  $x_j$   
 $-3 \leq x_j \leq 3$   
 $-0.3 \leq x_j \leq 0.3$  } acceptable ranges

- $0 \leq x_1 \leq 3$
- $-2 \leq x_2 \leq 0.5$
- $-100 \leq x_3 \leq 100$
- $-0.001 \leq x_4 \leq 0.001$
- $98.6 \leq x_5 \leq 105$

Which of these  
ranges need to be  
scaled?



# Some general hints with scaling

- Be careful when scaling the **time series** data! Why?
- To avoid **data leakage**, it is a good practice to fit the scaler on the training data and then use it to transform the testing data.
- Scaling the data does NOT change the shape of the **distributions**.





# Class Modules

- ✓ Module 1- Introduction to Deep Learning
- ✓ Module 2- Setting up Machine Learning Environment
- ✓ Module 3- Linear Regression (Econometrics approach)
- ✓ Module 4- Machine Learning Fundamentals
- Module 5- Linear Regression (Machine Learning approach)
- Module 6- Penalized Regression (Ridge, LASSO, Elastic Net)
- Module 7- Logistic Regression
- Module 8- K-Nearest Neighbors (KNN)
- Module 9- Classification and Regression Trees (CART)
- Module 10- Bagging and Boosting
- Module 11- Dimensionality Reduction (PCA)
- Module 12- Clustering (KMeans – Hierarchical)

