



Module 10- Part II

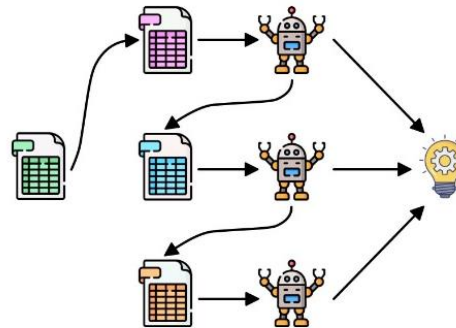
Boosting Models

AdaBoost, GBM, XGBoost



Prof. Pedram Jahangiry

Boosting

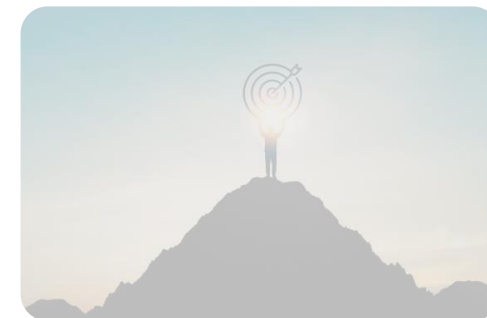


Sequential



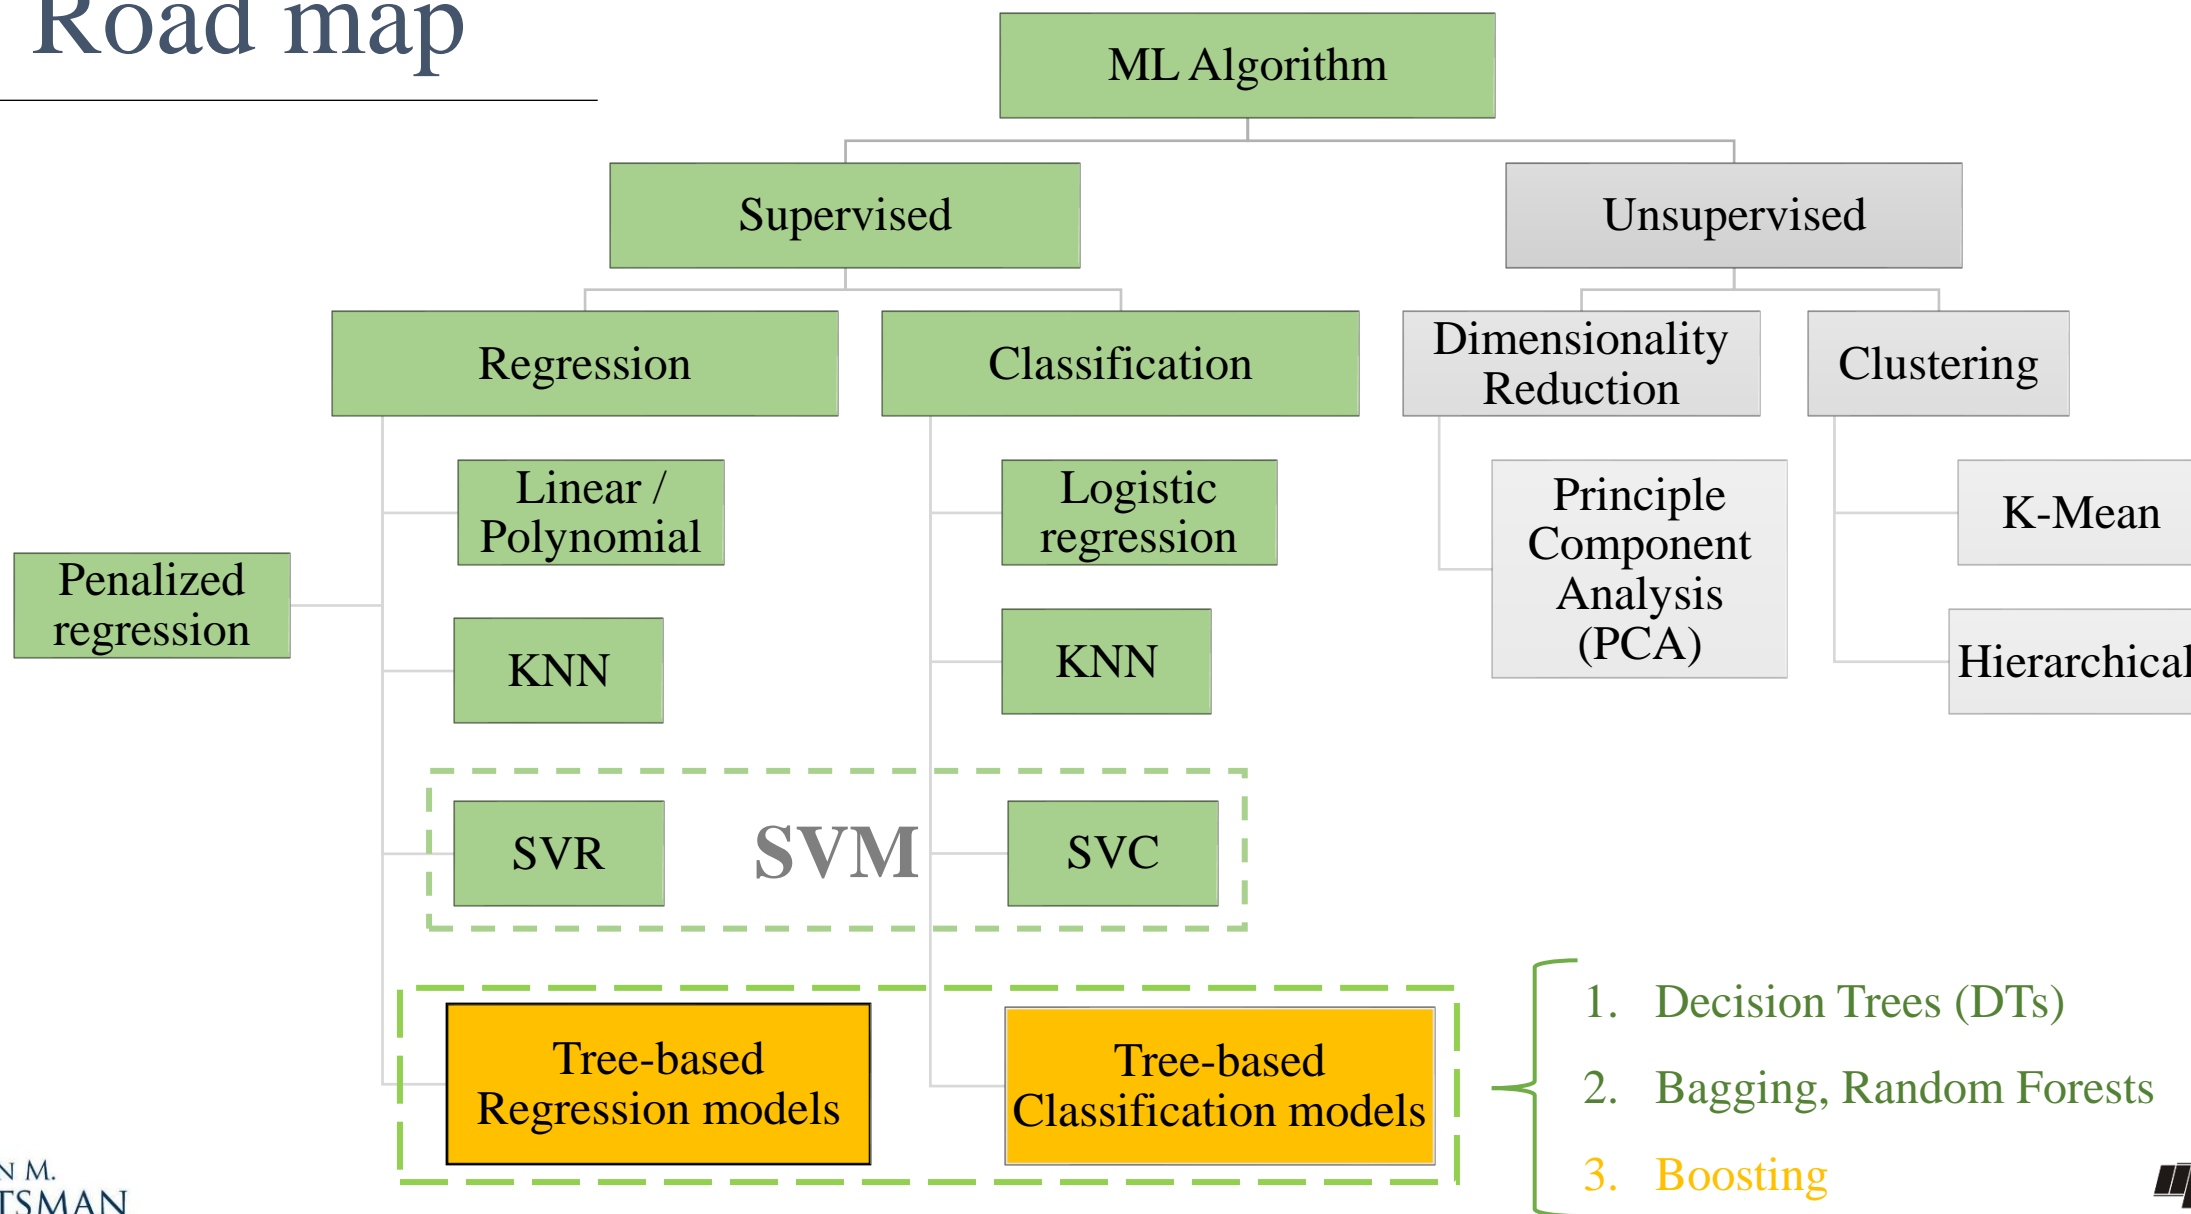
Class Modules

- Module 1- Introduction to Machine Learning
- Module 2- Setting up Machine Learning Environment
- Module 3- Linear Regression (Econometrics approach)
- Module 4- Machine Learning Fundamentals
- Module 5- Linear Regression (Machine Learning approach)
- Module 6- Penalized Regression (Ridge, LASSO, Elastic Net)
- Module 7- Logistic Regression
- Module 8- K-Nearest Neighbors (KNN)
- Module 9- Classification and Regression Trees (CART)
- **Module 10- Bagging and Boosting**
- Module 11- Dimensionality Reduction (PCA)
- Module 12- Clustering (KMeans – Hierarchical)





Road map





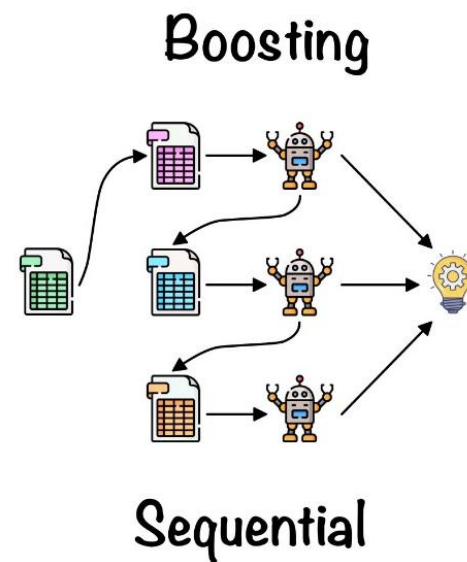
Topics

Part I

1. Bagging vs Boosting
2. AdaBoost
3. Gradient Boosting Machine (GBM)
4. XGBoost

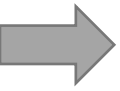
Part II

Pros and Cons



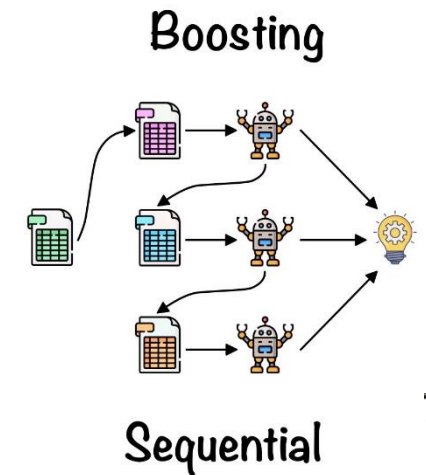
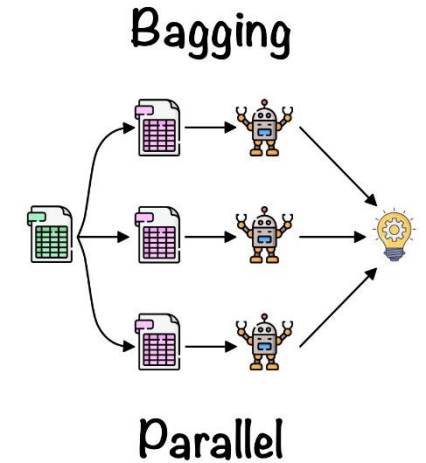
Part I

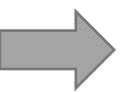
1. Bagging vs Boosting
2. AdaBoost
3. Gradient Boosting Machine (GBM)
4. XGBoost



Bagging vs Boosting

- **Bagging** consists of creating many “copies” of the training data (each copy is slightly different from another) and then apply the weak learner to each copy to obtain multiple weak models and then combine them.
- In bagging, the bootstrapped trees are **independent** from each other.
- **Boosting** consists of using the “original” training data and **iteratively** creating multiple models by using a weak learner. Each new model would be different from the previous ones in the sense that the weak learner, by building each new model tries to “fix” the **errors** which previous models make.
- In boosting, each tree is grown using information from **previous** tree.

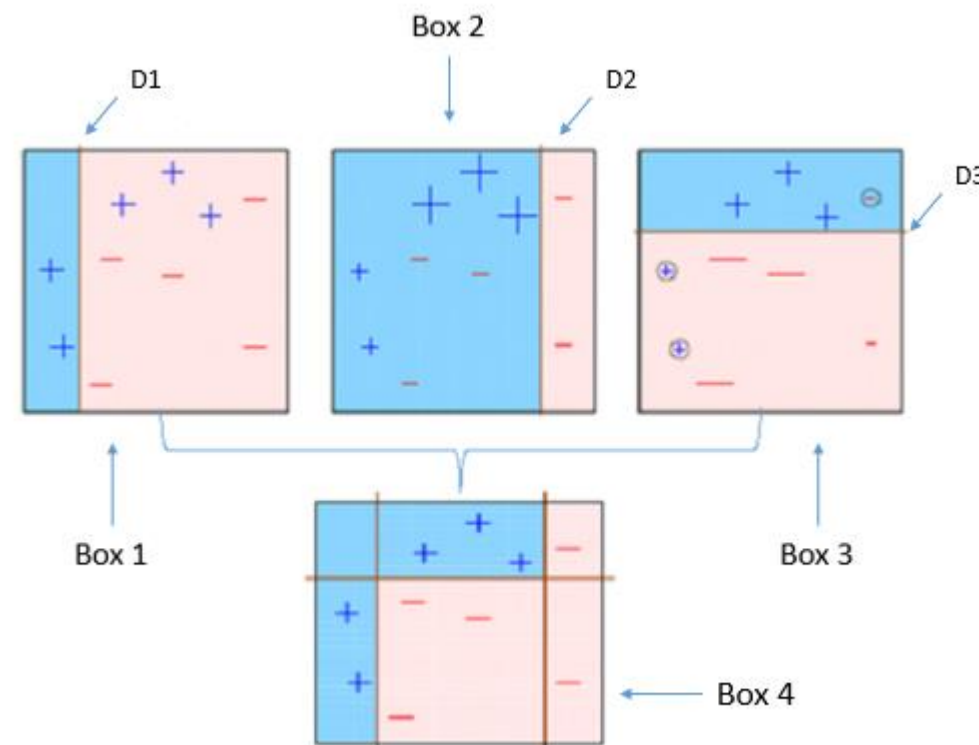




AdaBoost (Adaptive Boosting)

- Forest of weak learners (trees with only 1 feature; **stumps**).
- Each tree (stump) depends on the previous tree's **errors** rather than being independent.

- 1) Starting with usual splitting criteria!
- 2) Each tree (stump) gets **different weight** based on its prediction accuracy.
- 3) Each **observation** gets a weight inversely related to its predicted outcome. (ex, misclassified ones get more weight).
- 4) **Aggregation** is done based on each weak learner's weight.



Source: [Towards data science](#)



AdaBoost

Key features:

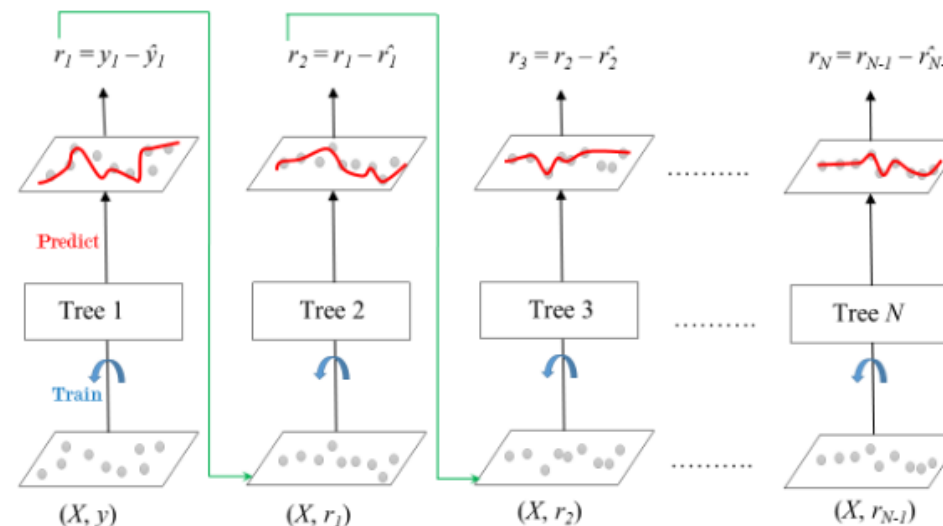
- Adaptive: Updates the weights of misclassified instances at each step.
- Tends to be sensitive to noise and outliers.
- Can be used with various base classifiers, but most commonly used with decision stumps.
- AdaBoost is old: AdaBoost is a popular boosting technique introduced by Yoav Freund and Robert Schapire in 1996.

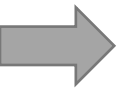


Gradient Boosting Machine (GBM)

Source: [Geeksforgeeks](https://www.geeksforgeeks.org/gradient-boosting/)

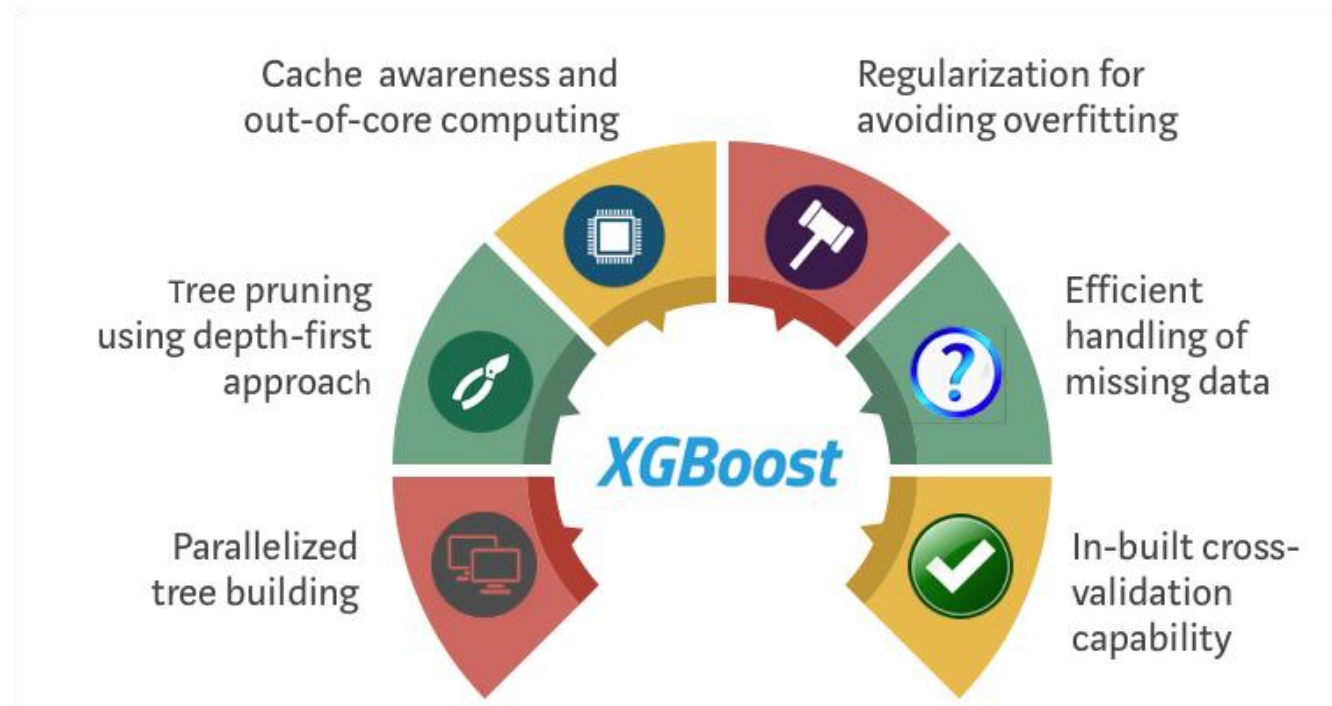
- In **gradient boosting**, each weak learner corrects its predecessor's error.
- Unlike AdaBoost, the **weights** of the training instances are not tweaked, instead, each predictor is trained using the **residual errors** of predecessor as labels.
- Unlike AdaBoost, each tree can be **larger than a stump**. However, the trees are still small. By fitting a small tree to the residuals, the GBM slowly improve \hat{f} in areas where it does not perform well.
- **Learning rate** shrinks the contribution of each tree. There is a trade-off between **learning rate** and **number of trees**. Learning rate slows down the process even further, allowing for more and different shaped trees to attack the residuals.
- **Aggregation** is done by adding the first tree predictions and a scaled (shrunk) version of the following trees.





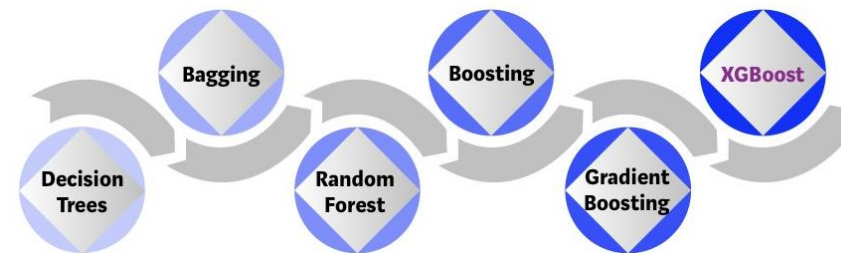
Extreme Gradient Boosting (XGBoost)

- **XGBoost** is a refined and customized version of a **gradient boosting** decision tree system, created with **performance** and **speed** in mind.
- **Extreme** refers to the fact that the algorithms and methods have been customized to push the limit of what is possible for gradient boosting algorithms.



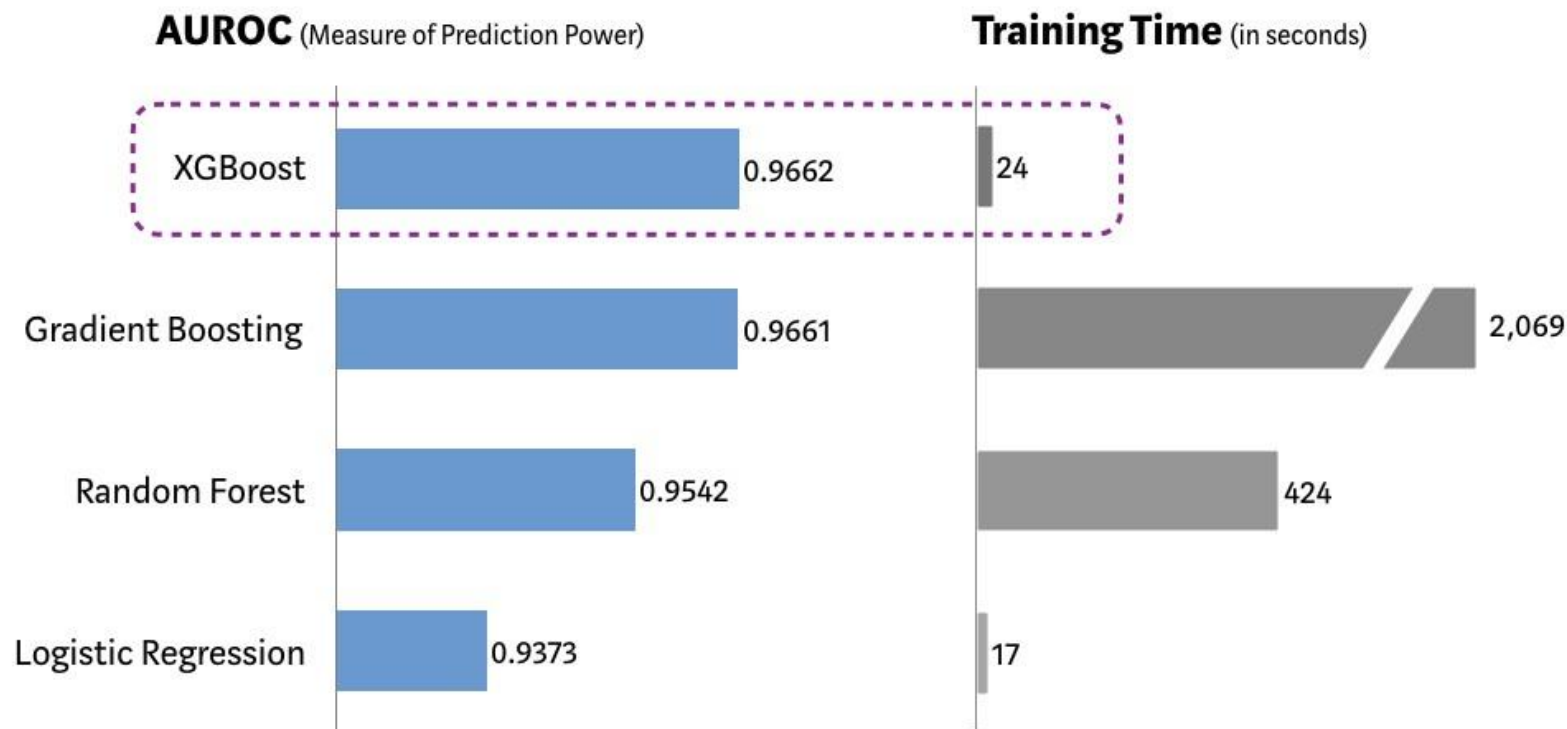


Put it all together!



Performance Comparison using SKLearn's 'Make_Classification' Dataset

(5 Fold Cross Validation, 1MM randomly generated data sample, 20 features)

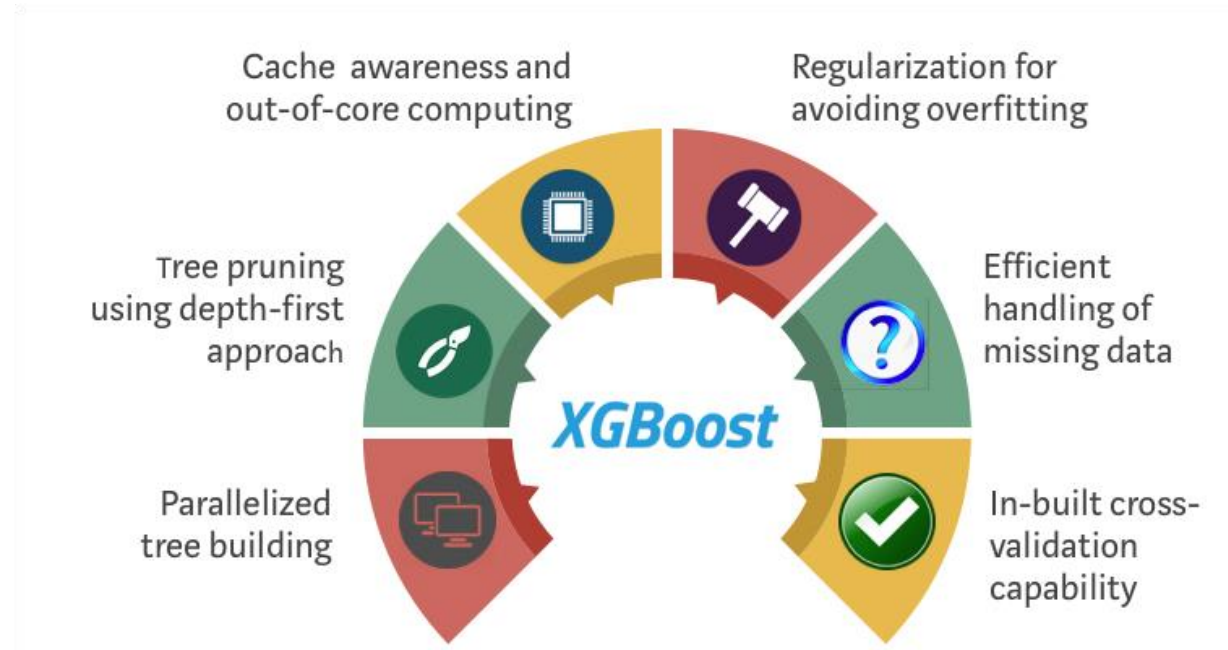


Part II

Pros and Cons

➔ XGBoost's Pros and Cons

Pros:



Cons:

- XGBoost is **more difficult to** understand, visualize and to **tune** compared to AdaBoost and random forests. There is a multitude of hyperparameters that can be tuned to increase performance.



Class Modules

- ✓ Module 1- Introduction to Machine Learning
- ✓ Module 2- Setting up Machine Learning Environment
- ✓ Module 3- Linear Regression (Econometrics approach)
- ✓ Module 4- Machine Learning Fundamentals
- ✓ Module 5- Linear Regression (Machine Learning approach)
- ✓ Module 6- Penalized Regression (Ridge, LASSO, Elastic Net)
- ✓ Module 7- Logistic Regression
- ✓ Module 8- K-Nearest Neighbors (KNN)
- ✓ Module 9- Classification and Regression Trees (CART)
- ✓ Module 10- Bagging and Boosting
- Module 11- Dimensionality Reduction (PCA)
- Module 12- Clustering (KMeans – Hierarchical)

